

# AI Impact Assessment System - Backend Strategy Document

**Version:** 1.0

**Date:** December 31, 2025

**Author:** Architecture Team

---

## Table of Contents

1. [Executive Summary](#)
  2. [System Architecture Overview](#)
  3. [Technology Stack](#)
  4. [Data Architecture](#)
  5. [API Design](#)
  6. [RAG Pipeline Design](#)
  7. [Multi-Agent LLM Orchestration](#)
  8. [Session & Audit Management](#)
  9. [Configuration Management](#)
  10. [Folder Structure](#)
  11. [Deployment Strategy](#)
  12. [Future Enhancements](#)
- 

## 1. Executive Summary


### 1.1 Purpose

This document outlines the backend strategy for an AI-powered Impact Assessment System that analyzes incoming project requirements and identifies impacted functional/technical modules to enable effective project planning and effort estimation.

## 1.2 Key Design Decisions

Decision	Choice	Rationale
Backend Framework	FastAPI	Modern, async, auto-generated OpenAPI docs
Vector Database	ChromaDB	Lightweight, local-first, good for POC scale
LLM Runtime	Ollama (local)	Privacy, no API costs, offline capability
LLM Model	phi3:mini	Lightweight, good structured output
Embedding Model	all-minilm	Very lightweight, sufficient for POC
Search Strategy	Hybrid (configurable)	Semantic + keyword matching

## 1.3 Core Workflow



## 2. System Architecture Overview

### 2.1 High-Level Architecture





## 2.2 Component Responsibilities

Component	Responsibility
API Gateway	Request routing, validation, error handling, CORS
Service Layer	Business logic, orchestration, data transformation
RAG Pipeline	Embedding generation, hybrid search, context retrieval
Agent Orchestrator	Multi-step LLM calls, prompt management, response parsing
Data Layer	Vector storage, file management, session persistence

## 3. Technology Stack

### 3.1 Core Technologies

TECHNOLOGY STACK

```
RUNTIME & FRAMEWORK
├── Python 3.11+
├── FastAPI (async web framework)
├── Uvicorn (ASGI server)
└── Pydantic v2 (data validation)

AI/ML COMPONENTS
├── Ollama (local LLM runtime)
│   ├── phi3:mini (text generation)
│   ├── all-minilm (embeddings)
│   ├── ChromaDB (vector database)
└── LangChain (optional - for complex chains)

DATA PROCESSING
├── Pandas (CSV/data manipulation)
├── python-multipart (file uploads)
└── aiofiles (async file I/O)

UTILITIES
├── python-dotenv (configuration)
├── structlog (structured logging)
└── httpx (async HTTP client for Ollama)
```

### 3.2 Version Requirements


Package	Version	Purpose
python	$\geq 3.11$	Runtime
fastapi	$\geq 0.109.0$	Web framework
chromadb	$\geq 0.4.22$	Vector database
ollama	$\geq 0.1.0$	LLM client
pandas	$\geq 2.1.0$	Data processing
pydantic	$\geq 2.5.0$	Validation

## 4. Data Architecture

### 4.1 Source Data Files

```
data/
└── raw/
    ├── epics.csv      # Master project initiatives
    ├── estimations.csv # Dev/QA effort with TDD embedded
    ├── tdds.csv       # Technical Design Documents
    ├── stories_tasks.csv # Jira-style work items
    └── gitlab_code.json # Code references
```

### 4.2 Data Relationships



### 4.3 ChromaDB Collections Schema

Collection: **epics**

Document: epic\_title + epic\_description (concatenated)

Metadata:

- epic\_id (string)
- epic\_status (string)
- epic\_priority (string)
- epic\_owner (string)
- epic\_team (string)
- epic\_start\_date (string)
- epic\_target\_date (string)

## Collection: estimations

Document: task\_description + tdd\_description (concatenated)

Metadata:

- estimation\_id (string)
- epic\_id (string)
- dev\_hours (float)
- qa\_hours (float)
- total\_hours (float)
- complexity (string)
- risk\_level (string)
- tdd\_technologies (list → JSON string)
- tdd\_dependencies (list → JSON string)

## Collection: tdds

Document: tdd\_section\_name + tdd\_section\_content\_summary (concatenated)

Metadata:

- tdd\_id (string)
- epic\_id (string)
- tdd\_section\_id (string)
- tdd\_title (string)
- tdd\_version (string)
- tdd\_status (string)
- tdd\_technologies (list → JSON string)
- tdd\_dependencies (list → JSON string)

## Collection: stories

Document: story\_title + acceptance\_criteria (concatenated)

Metadata:

- story\_id (string)
- epic\_id (string)
- tdd\_id (string)
- tdd\_section\_id (string)
- story\_type (string)
- story\_status (string)
- story\_points (int)
- story\_assignee (string)
- labels (list → JSON string)

**Collection:** `gitlab_code`

Document: code\_block\_description + functions\_defined (concatenated)

Metadata:

- code\_id (string)
- story\_id (string)
- epic\_id (string)
- tdd\_section\_id (string)
- gitlab\_repo (string)
- gitlab\_file\_path (string)
- code\_language (string)
- classes\_defined (list → JSON string)
- functions\_defined (list → JSON string)

## 5. API Design

### 5.1 API Endpoints Overview

API ENDPOINTS

BASE URL: /api/v1

## SESSION MANAGEMENT

| POST /session/create Create new assessment session | |

| GET /session/{session\_id} Get session details | |

| GET /session/{session\_id}/audit Get full audit trail | |

## STEP 1: INPUT REQUIREMENT

| POST /requirement/submit Submit requirement text | |

| POST /requirement/upload Upload supporting documents | |

| GET /requirement/{session\_id} Get submitted requirement | |

## STEP 2: HISTORICAL MATCHES

| POST /search/find-matches Trigger historical search | |

| POST /search/select-matches Confirm selected matches | |

| GET /search/{session\_id}/matches Get match results | |

## STEP 3: IMPACT RESULTS (5 separate agent calls)

POST	/impact/generate/modules	Generate impacted modules		
POST	/impact/generate/effort	Generate effort breakdown		
POST	/impact/generate/stories	Generate Jira stories		
POST	/impact/generate/code	Generate code impact analysis		
POST	/impact/generate/risks	Generate identified risks		
GET	/impact/{session_id}/summary	Get complete impact summary		

## CONFIGURATION

GET	/config	Get current configuration		
PUT	/config/search-weights	Update hybrid search weights		

## HEALTH & ADMIN

GET	/health	Health check		
POST	/admin/reindex	Reindex vector database		

## 5.2 Detailed API Specifications

### 5.2.1 Session Management

## **POST /api/v1/session/create**

```
json

// Request
{
  "user_id": "optional-user-identifier"
}

// Response
{
  "session_id": "sess_20251231_143052_abc123",
  "created_at": "2025-12-31T14:30:52Z",
  "status": "created",
  "audit_path": "sessions/2025-12-31/sess_20251231_143052_abc123/"
}
```

### **5.2.2 Step 1: Input Requirement**

## **POST /api/v1/requirement/submit**

```
json

// Request
{
  "session_id": "sess_20251231_143052_abc123",
  "requirement_description": "Implement OAuth 2.0 authentication with Azure AD B2C for member portal login with MFA support",
  "jira_epic_id": "EPIC-1234" // optional
}

// Response
{
  "session_id": "sess_20251231_143052_abc123",
  "requirement_id": "req_001",
  "status": "submitted",
  "character_count": 95,
  "extracted_keywords": ["OAuth 2.0", "Azure AD B2C", "authentication", "MFA", "member portal"]
}
```

## **POST /api/v1/requirement/upload**

```
json
```

```
// Request: multipart/form-data
// - session_id: string
// - files: File[] (.doc, .pdf, .png, .jpg)

// Response
{
  "session_id": "sess_20251231_143052_abc123",
  "uploaded_files": [
    {
      "file_id": "file_001",
      "filename": "requirements_spec.pdf",
      "size_bytes": 245000,
      "mime_type": "application/pdf",
      "stored_path": "uploads/sess_20251231_143052_abc123/requirements_spec.pdf"
    }
  ]
}
```

### 5.2.3 Step 2: Historical Matches

**POST /api/v1/search/find-matches**

json

```

// Request
{
  "session_id": "sess_20251231_143052_abc123",
  "search_config": {
    "semantic_weight": 0.7,    // configurable
    "keyword_weight": 0.3,    // configurable
    "max_results": 10
  }
}

// Response
{
  "session_id": "sess_20251231_143052_abc123",
  "matches": [
    {
      "match_id": "match_001",
      "epic_id": "EPIC-005",
      "epic_name": "Enterprise SSO Implementation",
      "description": "Implemented SAML 2.0 based SSO across all internal applications with MFA support",
      "match_score": 0.94,
      "score_breakdown": {
        "semantic_similarity": 0.92,
        "keyword_match": 0.88
      },
      "technologies": ["OAuth 2.0", "SAML", "Azure AD", "React", "Node.js"],
      "team_size": 5,
      "actual_hours": 480,
      "estimated_hours": 520,
      "variance": "-8%",
      "lessons_learned": "Early integration testing with identity provider saved significant debugging time"
    },
    // ... more matches
  ],
  "search_metadata": {
    "total_candidates_searched": 100,
    "search_duration_ms": 245,
    "weights_used": {
      "semantic": 0.7,
      "keyword": 0.3
    }
  }
}

```

## **POST /api/v1/search/select-matches**

```
json

// Request
{
  "session_id": "sess_20251231_143052_abc123",
  "selected_match_ids": ["match_001", "match_002", "match_004"]
}

// Response
{
  "session_id": "sess_20251231_143052_abc123",
  "selected_matches": 3,
  "status": "matches_confirmed",
  "ready_for_impact_analysis": true
}
```

### **5.2.4 Step 3: Impact Results (5 Agent Calls)**

#### **POST /api/v1/impact/generate/modules**

```
json
```

```

// Request
{
  "session_id": "sess_20251231_143052_abc123"
}

// Response
{
  "session_id": "sess_20251231_143052_abc123",
  "agent": "modules",
  "generated_at": "2025-12-31T14:35:22Z",
  "modules": {
    "functional": [
      {"name": "Member Login Flow", "impact": "HIGH", "reason": "Core authentication changes"},
      {"name": "Session Management", "impact": "HIGH", "reason": "Token-based session handling"},
      {"name": "Password Reset", "impact": "MEDIUM", "reason": "Integration with Azure AD B2C"},
      {"name": "Remember Me Feature", "impact": "LOW", "reason": "Minor token storage updates"}
    ],
    "technical": [
      {"name": "Auth Service", "impact": "HIGH", "reason": "New OAuth flow implementation"},
      {"name": "API Gateway", "impact": "HIGH", "reason": "Token validation middleware"},
      {"name": "Database Schema", "impact": "MEDIUM", "reason": "User token storage tables"},
      {"name": "Logging Service", "impact": "LOW", "reason": "Auth event logging"}
    ]
  },
  "total_modules": 10,
  "audit_file": "sessions/2025-12-31/sess_20251231_143052_abc123/agent_modules.json"
}

```

## POST /api/v1/impact/generate/effort

json

```

// Response
{
  "session_id": "sess_20251231_143052_abc123",
  "agent": "effort",
  "effort_breakdown": {
    "development": {
      "hours": 476,
      "breakdown": [
        {"module": "Auth Service", "hours": 160},
        {"module": "API Gateway", "hours": 120},
        // ...
      ]
    },
    "qa_testing": {
      "hours": 120,
      "breakdown": [...]
    },
    "design": {
      "hours": 40,
      "breakdown": [...]
    },
    "total_hours": 596
  },
  "story_points": 89,
  "historical_comparison": {
    "average": 420,
    "min": 280,
    "max": 520
  },
  "confidence": 0.85,
  "confidence_reasoning": "Based on 2 similar historical projects with 87%+ match"
}

```

## POST /api/v1/impact/generate/stories

json

```
// Response
{
  "session_id": "sess_20251231_143052_abc123",
  "agent": "stories",
  "stories": [
    {
      "story_id": "generated_001",
      "title": "Implement Azure AD B2C Integration",
      "type": "Story",
      "story_points": 13,
      "priority": "HIGH",
      "acceptance_criteria": [
        "User can authenticate via Azure AD B2C",
        "MFA is enforced for all logins",
        "Error handling for authentication failures"
      ],
      "linked_modules": ["Auth Service", "Member Login Flow"]
    },
    // ... more stories
  ],
  "total_stories": 10,
  "total_story_points": 89
}
```

## POST /api/v1/impact/generate/code

json

```

// Response
{
  "session_id": "sess_20251231_143052_abc123",
  "agent": "code_impact",
  "impacted_files": [
    {
      "file_path": "src/auth/AuthContext.tsx",
      "repository": "member-portal-frontend",
      "language": "TypeScript",
      "change_type": "modified",
      "reason": "Update auth context for OAuth flow"
    },
    {
      "file_path": "src/pages/Login.tsx",
      "repository": "member-portal-frontend",
      "language": "TypeScript",
      "change_type": "modified",
      "reason": "Integrate Azure AD B2C login button"
    },
    {
      "file_path": "src/auth/OAuthCallback.tsx",
      "repository": "member-portal-frontend",
      "language": "TypeScript",
      "change_type": "new",
      "reason": "Handle OAuth callback and token exchange"
    },
    // ... more files
  ],
  "total_files": 8,
  "summary": {
    "new_files": 2,
    "modified_files": 6,
    "repositories_affected": ["member-portal-frontend", "auth-service"]
  }
}

```

## POST /api/v1/impact/generate/risks

json

```
// Response
{
  "session_id": "sess_20251231_143052_abc123",
  "agent": "risks",
  "risks": [
    {
      "risk_id": "risk_001",
      "title": "Azure AD B2C Service Dependency",
      "severity": "HIGH",
      "description": "Authentication flow depends on Azure AD B2C availability. Service outage would prevent all logins.",
      "mitigation": "Implement fallback authentication mechanism and proper error handling"
    },
    {
      "risk_id": "risk_002",
      "title": "Token Migration Complexity",
      "severity": "MEDIUM",
      "description": "Existing users with active sessions need seamless migration to new token-based system.",
      "mitigation": "Implement grace period with dual-auth support during migration"
    },
    // ... more risks
  ],
  "total_risks": 5,
  "high_risks": 2,
  "medium_risks": 2,
  "low_risks": 1
}
```

## GET /api/v1/impact/{session\_id}/summary

json

```

// Response - Aggregated from all agent outputs
{
  "session_id": "sess_20251231_143052_abc123",
  "summary": {
    "total_effort_hours": 596,
    "story_points": 89,
    "modules_impacted": 10,
    "high_risks": 2
  },
  "estimation_confidence": {
    "score": 0.85,
    "based_on": "2 similar historical projects"
  },
  "modules": {...},
  "effort": {...},
  "stories": {...},
  "code_impact": {...},
  "risks": {...},
  "generated_at": "2025-12-31T14:40:15Z"
}

```

## 5.2.5 Configuration

### PUT /api/v1/config/search-weights

```

json


// Request
{
  "semantic_weight": 0.6,
  "keyword_weight": 0.4
}


// Response
{
  "status": "updated",
  "weights": {
    "semantic_weight": 0.6,
    "keyword_weight": 0.4
  },
  "validation": "Weights sum to 1.0 ✓"
}

```

## 6. RAG Pipeline Design

### 6.1 Hybrid Search Architecture





## 6.2 Embedding Strategy

### Document Chunking & Embedding:

Collection	Fields to Embed	Chunking Strategy
epics	epic_title + epic_description	Single chunk per epic
estimations	task_description + tdd_description	Single chunk per estimation
tdd	tdd_section_name + tdd_section_content_summary	Single chunk per TDD section
stories	story_title + acceptance_criteria	Single chunk per story
gitlab_code	code_block_description + functions_defined	Single chunk per code block

### Embedding Generation:

## EMBEDDING FLOW (using all-minilm)

Raw Text —► Preprocessing —► Ollama API —► Vector (384-dim)

Preprocessing:

1. Lowercase
2. Remove extra whitespace
3. Truncate to max 512 tokens

## 7. Multi-Agent LLM Orchestration

### 7.1 Agent Architecture

#### AGENT ORCHESTRATION FLOW


User Triggers "Generate Impact Results"



#### ORCHESTRATOR

- Loads session context (requirement + selected matches)
- Sequences agent calls
- Manages audit logging
- Handles errors and retries





## 7.2 Agent Specifications

### Agent 1: Modules Agent

**Purpose:** Identify impacted functional and technical modules

#### Input Context:

- Requirement description
- Selected historical matches

#### Prompt Template:

You are an expert software architect analyzing project requirements.

**REQUIREMENT:**

{requirement\_description}

**SIMILAR HISTORICAL PROJECTS:**

{formatted\_historical\_matches}

Based on this requirement and historical data, identify the impacted modules.

**OUTPUT FORMAT (JSON):**

```
{  
  "functional_modules": [  
    {"name": "string", "impact": "HIGH|MEDIUM|LOW", "reason": "string"}  
  ],  
  "technical_modules": [  
    {"name": "string", "impact": "HIGH|MEDIUM|LOW", "reason": "string"}  
  ]  
}
```

Provide exactly 10 modules total (mix of functional and technical).

**Output Schema:**

```
json  
  
{  
  "functional_modules": [  
    {"name": "string", "impact": "enum", "reason": "string"}  
  ],  
  "technical_modules": [  
    {"name": "string", "impact": "enum", "reason": "string"}  
  ]  
}
```

---

**Agent 2: Effort Agent**

**Purpose:** Generate effort breakdown by activity type

**Input Context:**

- Requirement description
- Selected historical matches (with actual/estimated hours)
- Generated modules (from Agent 1)

### Prompt Template:

You are a project estimation expert.

#### REQUIREMENT:

{requirement\_description}

#### HISTORICAL EFFORT DATA:

{historical\_effort\_data}

Provide effort estimates based on historical patterns.

#### OUTPUT FORMAT (JSON):

```
{  
  "development_hours": number,  
  "qa_testing_hours": number,  
  "design_hours": number,  
  "total_hours": number,  
  "story_points": number,  
  "confidence_score": 0.0-1.0,  
  "confidence_reasoning": "string"  
}
```

### Agent 3: Stories Agent

**Purpose:** Generate Jira stories with acceptance criteria

### Input Context:

- Requirement description
- Generated modules (from Agent 1)
- Historical stories patterns

### Prompt Template:

You are a product owner creating user stories.

**REQUIREMENT:**

{requirement\_description}

**IMPACTED MODULES:**

{modules}

Generate user stories for implementation.

**OUTPUT FORMAT (JSON):**

```
{
  "stories": [
    {
      "title": "string",
      "type": "Story|Task",
      "story_points": number,
      "priority": "HIGH|MEDIUM|LOW",
      "acceptance_criteria": ["string"],
      "linked_modules": ["string"]
    }
  ]
}
```

Generate exactly 10 stories.

---

## Agent 4: Code Impact Agent

**Purpose:** Identify files and repositories that will be impacted

**Input Context:**

- Requirement description
- Generated modules (from Agent 1)
- Historical code patterns from similar projects

**Prompt Template:**

You are a senior developer analyzing code impact.

**REQUIREMENT:**

{requirement\_description}

**IMPACTED MODULES:**

{modules}

**HISTORICAL CODE PATTERNS:**

{historical\_code}

Identify files and repositories that will be affected.

**OUTPUT FORMAT (JSON):**

```
{
  "impacted_files": [
    {
      "file_path": "string",
      "repository": "string",
      "language": "string",
      "change_type": "new|modified",
      "reason": "string"
    }
  ]
}
```

Generate 8-12 impacted files.

## Agent 5: Risks Agent

**Purpose:** Identify potential risks and mitigation strategies

**Input Context:**

- Requirement description
- Generated modules (from Agent 1)
- Historical lessons learned

**Prompt Template:**

You are a risk analyst reviewing project requirements.

**REQUIREMENT:**

{requirement\_description}

**IMPACTED MODULES:**

{modules}

**HISTORICAL LESSONS LEARNED:**

{lessons\_learned}

Identify potential risks and mitigation strategies.

**OUTPUT FORMAT (JSON):**

```
{
  "risks": [
    {
      "title": "string",
      "severity": "HIGH|MEDIUM|LOW",
      "description": "string",
      "mitigation": "string"
    }
  ]
}
```

Generate 5 risks with varying severity levels.

### 7.3 LLM Configuration

OLLAMA CONFIGURATION

Model: phi3:mini

Parameters:


- └── temperature: 0.3 (lower for more consistent output)
- └── top\_p: 0.9
- └── max\_tokens: 2048
- └── format: json (for structured output)

System Message:

"You are a helpful AI assistant that always responds in valid JSON format. Follow the output schema exactly."

## 8. Session & Audit Management

### 8.1 Session Lifecycle



### 8.2 Audit Folder Structure

```
sessions/
└── 2025-12-31/
    └── sess_20251231_143052_abc123/
        ├── session_metadata.json      # Session info and status
        ├── step1_input/
        │   ├── requirement.json      # Submitted requirement text
        │   ├── extracted_keywords.json # AI-extracted keywords
        │   └── uploads/               # Uploaded documents
```

```
└── requirements_spec.pdf  
└── architecture_diagram.png  
  
└── step2_search/  
    ├── search_request.json      # Search configuration used  
    ├── all_matches.json         # All returned matches  
    └── selected_matches.json   # User-selected matches  
  
└── step3_agents/  
    ├── agent_modules.json      # Modules agent output  
        ├── input_prompt.txt    # Full prompt sent  
        ├── raw_response.txt   # Raw LLM response  
        └── parsed_output.json  # Parsed structured output  
  
    ├── agent_effort.json       # Effort agent output  
    │   ...  
    ├── agent_stories.json      # Stories agent output  
    │   ...  
    ├── agent_code.json         # Code impact agent output  
    │   ...  
    └── agent_risks.json        # Risks agent output  
        ...  
  
└── final_summary.json      # Aggregated final output
```

### 8.3 Session Metadata Schema

json

```
{  
  "session_id": "sess_20251231_143052_abc123",  
  "created_at": "2025-12-31T14:30:52Z",  
  "updated_at": "2025-12-31T14:45:30Z",  
  "status": "completed",  
  "user_id": "optional-user-id",  
  "steps_completed": {  
    "step1_requirement": true,  
    "step2_matches": true,  
    "step3_modules": true,  
    "step3_effort": true,  
    "step3_stories": true,  
    "step3_code": true,  
    "step3_risks": true  
  },  
  "timing": {  
    "step1_duration_ms": 1200,  
    "step2_search_duration_ms": 2450,  
    "step3_modules_duration_ms": 3200,  
    "step3_effort_duration_ms": 2800,  
    "step3_stories_duration_ms": 4100,  
    "step3_code_duration_ms": 3500,  
    "step3_risks_duration_ms": 2900,  
    "total_duration_ms": 20150  
  },  
  "configuration_used": {  
    "search_weights": {  
      "semantic": 0.7,  
      "keyword": 0.3  
    },  
    "llm_model": "phi3:mini",  
    "embedding_model": "all-minilm"  
  }  
}
```

## 9. Configuration Management

### 9.1 Configuration Schema

yaml

```
# config/settings.yaml

app:
  name: "AI Impact Assessment System"
  version: "1.0.0"
  environment: "development" # development / production

server:
  host: "0.0.0.0"
  port: 8000
  cors_origins:
    - "http://localhost:3000"
    - "http://localhost:5173"

ollama:
  base_url: "http://localhost:11434"
  generation_model: "phi3:mini"
  embedding_model: "all-minilm"
  timeout_seconds: 120
  generation_params:
    temperature: 0.3
    top_p: 0.9
    max_tokens: 2048

chromadb:
  persist_directory: "./data/chroma"
  collection_prefix: "impact_assessment"
  embedding_dimension: 384

search:
  default_weights:
    semantic: 0.70
    keyword: 0.30
  max_results: 10
  min_score_threshold: 0.3

data:
  raw_data_path: "./data/raw"
  uploads_path: "./data/uploads"
  sessions_path: "./data/sessions"

logging:
  level: "INFO"
```

```
format: "json"
file_path: "./logs/app.log"
```

## 9.2 Environment Variables

```
bash

# .env file

# Application
APP_ENV=development
DEBUG=true

# Ollama
OLLAMA_BASE_URL=http://localhost:11434
OLLAMA_GEN_MODEL=phi3:mini
OLLAMA_EMBED_MODEL=all-minilm

# ChromaDB
CHROMA_PERSIST_DIR=./data/chroma

# Search Weights (defaults, configurable via API)
SEARCH_WEIGHT_SEMANTIC=0.70
SEARCH_WEIGHT_KEYWORD=0.30

# Paths
DATA_RAW_PATH=./data/raw
DATA_UPLOADS_PATH=./data/uploads
DATA_SESSIONS_PATH=./data/sessions
```

## 10. Folder Structure

```
ai-impact-assessment/
|
|   app/
|   |   __init__.py
|   |   main.py          # FastAPI application entry point
|   |
|   |   api/
|   |   |   __init__.py
```

```
|- routes/
  |- __init__.py
  |- session.py      # Session management endpoints
  |- requirement.py # Step 1 endpoints
  |- search.py       # Step 2 endpoints
  |- impact.py       # Step 3 endpoints
  |- config.py       # Configuration endpoints
  |- health.py       # Health check endpoints

  |- dependencies.py # FastAPI dependencies

|- core/
  |- __init__.py
  |- config.py       # Configuration management
  |- logging.py      # Logging configuration

|- models/
  |- __init__.py
  |- request.py      # Pydantic request models
  |- response.py     # Pydantic response models
  |- domain.py       # Domain models

|- services/
  |- __init__.py
  |- session_service.py # Session management logic
  |- requirement_service.py # Requirement processing
  |- search_service.py   # Hybrid search logic
  |- impact_service.py  # Impact analysis orchestration
  |- file_service.py    # File upload handling

|- agents/
  |- __init__.py
  |- base_agent.py    # Base agent class
  |- modules_agent.py # Modules identification agent
  |- effort_agent.py  # Effort estimation agent
  |- stories_agent.py # Story generation agent
  |- code_agent.py    # Code impact agent
  |- risks_agent.py   # Risk identification agent
  |- orchestrator.py  # Agent orchestration

|- rag/
  |- __init__.py
  |- embeddings.py    # Embedding generation
  |- vector_store.py  # ChromaDB operations
```

```
    └── hybrid_search.py      # Hybrid search implementation

    └── utils/
        ├── __init__.py
        ├── audit.py          # Audit logging utilities
        └── helpers.py         # General utilities

    └── data/
        ├── raw/              # Source CSV/JSON files
            ├── epics.csv
            ├── estimations.csv
            ├── tdds.csv
            ├── stories_tasks.csv
            └── gitlab_code.json

        ├── chroma/           # ChromaDB persistence

        ├── uploads/           # Uploaded documents

        └── sessions/          # Session audit trails
            └── {date}/
                └── {session_id}/

    └── config/
        └── settings.yaml     # Main configuration

    └── scripts/
        ├── init_vector_db.py # Initialize ChromaDB with data
        ├── test_ollama.py    # Test Ollama connection
        └── reindex.py         # Reindex vector database

    └── tests/
        ├── __init__.py
        ├── test_api/
        ├── test_services/
        ├── test_agents/
        └── test_rag/

    └── logs/               # Application logs

    └── docs/               # Documentation
        └── DATA_SCHEMA_README.md

    └── .env                 # Environment variables
```

```
└── .env.example           # Example environment file
    ├── requirements.txt    # Python dependencies
    ├── pyproject.toml       # Project metadata
    ├── Makefile             # Common commands
    └── README.md            # Project README
```

## 11. Deployment Strategy

### 11.1 Local Development (MacBook)

#### LOCAL DEVELOPMENT SETUP

##### Prerequisites:

1. Python 3.11+ installed
2. Ollama installed (brew install ollama)
3. Pull required models:
  - ollama pull phi3:mini
  - ollama pull all-minilm

##### Startup Sequence:

```
Terminal 1: Start Ollama
$ ollama serve


Terminal 2: Start FastAPI
$ cd ai-impact-assessment
$ python -m venv venv
$ source venv/bin/activate
$ pip install -r requirements.txt
$ python scripts/init_vector_db.py # First time only
$ uvicorn app.main:app --reload --port 8000

Terminal 3: Start Frontend (if applicable)
$ cd frontend && npm run dev
```

#### Access Points:

- API: <http://localhost:8000>
- API Docs: <http://localhost:8000/docs>
- Ollama: <http://localhost:11434>

## 11.2 AWS Production Architecture (Future)



## 12. Future Enhancements

### 12.1 Phase 2 Enhancements

Feature	Description	Priority
Real-time Jira Integration	Create tickets directly in Jira	High
Document Content Extraction	Parse PDF/DOC content for context	High
Streaming Responses	Stream agent outputs for better UX	Medium
User Authentication	Add user management and auth	Medium
Export to PDF/Excel	Generate downloadable reports	Medium

### 12.2 Phase 3 Enhancements

Feature	Description	Priority
GitLab Integration	Pull real code structure	High
Feedback Loop	Learn from actual vs estimated	High
Multi-tenant Support	Support multiple organizations	Medium
A/B Testing for Prompts	Optimize agent prompts	Low
Dashboard Analytics	Historical accuracy metrics	Low

### 12.3 Scalability Considerations

SCALING DECISION MATRIX

Data Volume	Current	Recommended Action
< 1,000 records	ChromaDB	Keep current setup
1,000-10,000	ChromaDB	Add indexing, pagination
10,000-100,000	Pinecone	Migrate to managed vector DB
> 100,000	OpenSearch	Enterprise vector search
Concurrent Users	Current	Recommended Action
< 10 users	Single server	Keep current setup
10-50 users	Single server	Add caching (Redis)
50-200 users	Load balanced	Add multiple instances
> 200 users	ECS/K8s	Full containerized deploy

## Appendix A: API Error Codes

Code	HTTP Status	Description
SESSION_NOT_FOUND	404	Session ID does not exist
INVALID_SESSION_STATE	400	Operation not allowed in current session state
REQUIREMENT_TOO_SHORT	400	Requirement text < 20 characters
FILE_TYPE_NOT_ALLOWED	400	Uploaded file type not supported
SEARCH_WEIGHTS_INVALID	400	Search weights don't sum to 1.0
OLLAMA_UNAVAILABLE	503	Cannot connect to Ollama server
LLM_TIMEOUT	504	LLM response exceeded timeout
VECTOR_DB_ERROR	500	ChromaDB operation failed

## Document End

*This strategy document serves as the foundation for implementing the AI Impact Assessment System backend. All implementation should follow the patterns and structures defined herein.*

